

Introduction to OpenMP



ICHEC
Irish Centre for High-End Computing



Dr. Prithwish Nandi
Dr. Adam Ralph
UCD, June 2018



An Roinn Post, Fiontar agus Nuálaíocht
Department of Jobs, Enterprise and Innovation



AN tEolas
DOIRTEACHAIS AGUS INGLANNA
DEPARTMENT OF
EDUCATION AND SKILLS

HEA
Higher Education Authority
an tArd-Chomhairle na n-Éireann

Outline

- Introduction to OpenMP
- OpenMP Directives
 - Directive format
 - Parallel Construct
 - Work-Sharing Constructs
 - Synchronisation Constructs
- Performance Considerations
- OpenMP versions and features

What is OpenMP?


- Open Multi-Processing.
- API to explicitly specify multi-threaded, shared-memory parallelism.
- Three primary API components:
 - Compiler directives
 - Runtime library routines
 - Environment variables
- Designed for C, C++ and Fortran.
- Open standard for portable and scalable parallel programming, multi-platform.

Terminology


- **OpenMP thread:** a lightweight process - an instance of the programme and data.
- **thread team:** a set of threads co-operating on a task.
- **master thread:** the co-ordinating thread.
- **thread-safety:** threads are not interfering with data from other threads.
- **OpenMP directive:** pre-processed OpenMP code.
- **construct:** an OpenMP executable directive.
- **clause:** controls the operation of the directive or data.

OpenMP Programming Model ... (1)

- Multi-processor/core systems.
- Shared-memory UMA or NUMA systems.



Uniform Memory Access




Non-Uniform Memory Access

OpenMP Programming Model ... (2)

- Thread Based Parallelism
 - Execute a block of code with multiple threads.
 - Thread is the smallest unit of processing scheduled by an OS.
 - Typically, number of threads = number of cores.
 - But, can be decided by the application and/or programmer.
- Explicit Parallelism
 - OpenMP directives, library functions and environment variables explicitly specify what and how to parallelise a block of code.

OpenMP Programming Model ... (3)


- Fork-Join Model



- Nested parallelism
 - Placement of parallel regions inside parallel regions
- Dynamic threads
 - Runtime environment can dynamically alter number of threads used to execute a parallel region.

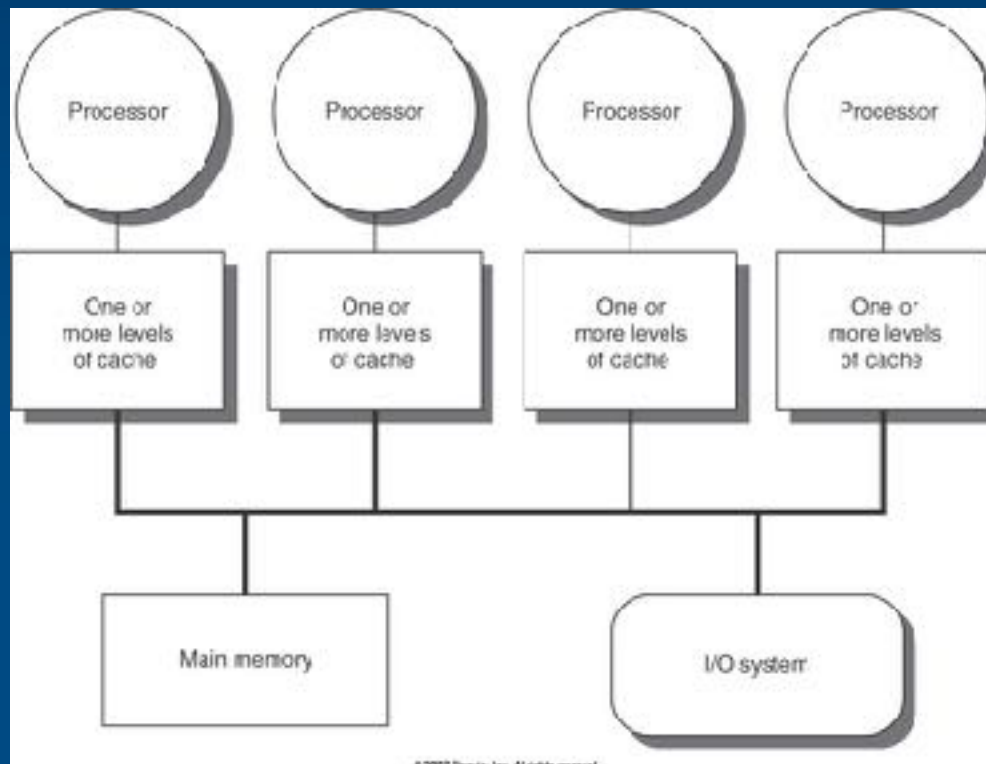
Memory model ... (1)

- Threads have access to shared data.
- Threads communicate by sharing data (variables) with other threads.
- Threads have private data.



Memory model ... (2)

- Threads allowed to have its own temporary view of memory (cache, registers) to avoid going to memory for every reference to a variable.
- If multiple threads write the same memory unit then a data race occurs.
 - Use synchronization to protect data conflicts.
 - Careless use of synchronization can lead to dead-locks



© 2007 Oracle, Inc. All rights reserved.

OpenMP Components

- Compiler Directives
 - Specify a parallel region
 - Divide work among threads
 - Synchronize threads
- Runtime Library Functions
 - Set and query thread-related information
 - number of threads, identifier, team size, ...
 - Query if in parallel region
 - Implement and coordinate locks
 - Query timing-related information
 - wall-clock time and resolution
- Environment Variables
 - Specify number of threads
 - Divide work among threads
 - Bind threads to processors

Writing an OpenMP Program

- Decompose the problem into tasks
 - Ideally, these tasks can be worked on independently of the others.
- Map tasks onto “threads of execution” (processors/cores)
- Decide data type. Threads have *shared* and *private* data
 - Shared: used by more than one thread
 - Private: local to each thread
- Write source code using compiler directives and library functions.
- Decide the necessary environment variables.
- Choices may depend on (among many things)
 - The nature of the problem
 - The level of performance needed



Compiling:

	Compiler	Flag
Intel	icc (C) icpc (C++) ifort (Fortran)	-qopenmp -openmp
GNU	gcc (C) g++ (C++) g77/gfortran (Fortran)	-fopenmp

See: <http://openmp.org/wp/openmp-compilers/> for the full list.

Outline

- Introduction to OpenMP
- OpenMP Directives
 - Directive format
 - Parallel Construct
 - Work-Sharing Constructs
 - Synchronization Constructs
- Performance Considerations
- OpenMP versions and features

Directive Format

C/C++:

```
#pragma omp directive-name [clause[clause]...]
{
    block of code
}
```

Fortran free form:

```
!$omp directive-name [clause[clause]...]
    block of code
!$omp end directive-name
```

Fortran fixed form:

```
!$omp | c$omp | *$omp directive-name [clause[clause]...]
    block of code
!$omp | c$omp | *$omp end directive-name
```

Outline

- Introduction to OpenMP
- OpenMP Directives
 - Directive format
 - Parallel Construct
 - Work-Sharing Constructs
 - Synchronization Constructs
- Performance Considerations
- OpenMP versions and features

Parallel Construct

- The fundamental construct in OpenMP.
- Creates team of threads
- Every thread executes the same statements inside the parallel region at the end of the parallel region there is an implicit barrier

C/C++:

```
#pragma omp parallel [clauses]  
{  
    ...  
}
```

Fortran:

```
!$omp parallel [clauses]  
...  
!$omp end parallel
```


Parallel Construct

- Clauses:

```
num_threads (integer-expression)  
if (scalar-expression)
```

- Data Clauses:

```
private (list)  
shared (list)  
default (shared | none)  
firstprivate (list)  
reduction (operator: list)  
copyin (list)
```

How many threads?

The number of threads in a parallel region is determined by:

- Use of `num_threads(n)` clause.
- Setting of the `OMP_NUM_THREADS` environment variable.
- Use of the `omp_set_num_threads(n)` library function.
- The implementation default - usually the number of CPUs/cores on a node

Threads are numbered from 0 (master thread) to $n-1$ where n =the total number of threads.

Example

```
void main()  
{  
    double Res[1000];  
    #pragma omp parallel num_threads(4)  
    {  
        block of code  
    }  
}
```

- Environment Variables

- To control the execution of parallel program at run-time.
- csh/tcsh: `setenv OMP_NUM_THREADS n`
- ksh/sh/bash: `export OMP_NUM_THREADS=n`
`echo $OMP_NUM_THREADS`

- Runtime Functions

- To manage the parallel program dynamically.
- `omp_set_num_threads(n)` - set the desired number of threads
- `omp_get_num_threads()` - returns the current number of threads
- `omp_get_thread_num()` - returns the id of this thread
- `omp_in_parallel()` - returns `.true.` if inside parallel region

C/C++: Add `#include<omp.h>`

Fortran: Add use `omp_lib`

```
double A[1000];
```

```
omp_set_num_threads(4);
```

```
foo(0,A); foo(1,A); foo(2,A); foo(3,A);
```

```
printf("All Done\n");
```

```
double A[1000];
omp_set_num_threads(4);
#pragma omp parallel
{
    int tid = omp_get_thread_num();
    foo(tid,A);
}
printf("All Done\n");
```

Hello World

C - Serial:

```
#include<stdio.h>

int main(int argc, char**argv)
{
    printf("Hello world!\n");
}
```

C - Parallel:

```
#include<stdio.h>
#include<omp.h>
int main(int argc, char**argv)
{
    #pragma omp parallel
    printf("Hello from thread %d out of %d\n", omp_get_thread_num(),
                                                omp_get_num_threads());
}
```

Hello World

Fortran - Serial:

```
program hello
implicit none

print *, 'Hello world!'

end program hello
```

Fortran - Parallel:


```
program hello
use omp_lib
implicit none

!$omp parallel
print *, 'Hello from thread', omp_get_thread_num(), 'out of', &
    omp_get_num_threads()
!$omp end parallel

end program hello
```

Nested parallel regions:

- If a parallel directive is encountered within another parallel directive, a new team of threads will be created.
- `omp_set_nested()`, `OMP_NESTED`, `omp_get_nested()`
- Num threads affects the new regions
- New threads with one thread unless nested parallelism is enabled
- `num_threads(n)` clause or dynamic threading for different num threads



Dynamic threads:

- Used to create a parallel region with a variable number of threads
- `omp_set_dynamic()`, `OMP_DYNAMIC`, `omp_get_dynamic()`
- OpenMP runtime will decide the number of threads

```
omp_set_dynamic(0);  
omp_set_num_threads(10);  
#pragma omp parallel  
printf("Num threads in non-dynamic region = %d\n", omp_get_num_threads());  
  
omp_set_dynamic(1);  
omp_set_num_threads(10);  
#pragma omp parallel  
printf("Num threads in dynamic region is = %d\n", omp_get_num_threads());
```

Parallel Construct

- Clauses:

`num_threads (integer-expression)`

`if (scalar-expression)`

`nowait`

- Data Clauses:

`private (list)`

`shared (list)`

`default (shared | none)`

`firstprivate (list)`

`reduction (operator: list)`

`copyin (list)`

If Clause:

- Used to make the parallel region directive itself conditional.
- Only execute in parallel if expression is true.

C/C++:

```
#pragma omp parallel if(n>100)
{
    ...
}
```

Fortran:

```
!$omp parallel if(n>100)
    ...
!$omp end parallel
```

nowait Clause:

- allows threads that finish earlier to proceed without waiting

```
#pragma omp parallel nowait
{
    ...
}
```

```
!$omp parallel
    ...
!$omp end parallel nowait
```

Parallel Construct

- Clauses:

```
num_threads (integer-expression)  
if (scalar_expression)  
nowait
```

- Data Clauses:

```
private (list)  
shared (list)  
default (shared | none)  
firstprivate (list)  
reduction (operator: list)  
copyin (list)
```

Data Clauses

- Used in conjunction with several directives to control the scoping of enclosed variables.
 - `default(shared|none)`: The default scope for all of the variables; private for Fortran
 - `shared(list)`: Variable is shared by all threads in the team. All threads can read or write to that variable.
 - C/C++: `#pragma omp parallel default(none) shared(n)`
 - Fortran: `!$omp parallel default(none) shared(n)`
 - `private(list)`: Each thread has a private copy of variable. It can only be read or written by its own thread.
 - C/C++: `#pragma omp parallel default(shared) private(tid)`
 - Fortran: `!$omp parallel default(shared) private(tid)`

Example

C:

```
#include<stdio.h>
#include<omp.h>
int main()
{
    int tid, nthreads;
    #pragma omp parallel private(tid), shared(nthreads)
    {
        tid = omp_get_thread_num();
        nthreads = omp_get_num_threads();
        printf("Hello from thread %d out of %d\n", tid, nthreads);
    }
}
```

Example

Fortran:

```

program hello
use omp_lib
implicit none
integer tid, nthreads

!$omp parallel private(tid), shared(nthreads)
tid = omp_get_thread_num()
nthreads = omp_get_num_threads()
print*, 'Hello from thread',tid,'out of',nthreads
!$omp end parallel

end program hello
  
```

- How do we decide which variables should be shared and which private?
 - Loop indices - private
 - Loop temporaries - private
 - Read-only variables - shared
 - Main arrays - shared
- Most variables are shared by default
 - C/C++: File scope, static variables
 - Fortran: COMMON blocks, SAVE, MODULE variables
 - Both: dynamically allocated variables
- Variables declared in parallel region are always private

Additional Data Clauses

- **firstprivate(list)**: pre-initialize private vars with value of variable with same name before parallel construct.
- **lastprivate(list)**: On exiting the parallel region, this gives private data the value of last iteration (if sequential)
- **threadprivate(list)**: Used to make global file scope variables (C/C++) or common blocks (Fortran) private to thread.
- **copyin(list)**: Copies a value from master thread to all threadprivate variables of a thread team.

```
j = jstart;
#pragma omp parallel for firstprivate(j)
{
    for(i=1; i<=n; i++)
    {
        j = j + 1;
        a[i] = a[i] + j;
    }
}
```

```
#pragma omp parallel for lastprivate(x)
{
    for(i=1; i<=n; i++)
    {
        x = sin( pi * dx * (float)i );
        a[i] = exp(x);
    }
}
lastx = x;
```

Outline

- Introduction to OpenMP
- OpenMP Directives
 - Directive format
 - Parallel Construct
 - Work-Sharing Constructs
 - Synchronization Constructs
- Performance Considerations
- OpenMP versions and features

Work-Sharing Constructs

- Work-sharing construct divides work among the thread team; what kind of work to be parallelised.
- Specify inside the parallel region.
- No new threads created. Construct must be encountered by all threads in the team.
- No implied barrier on entry to a work-sharing construct; Yes at end of construct.
- Work-sharing Constructs:
 - Loop
 - Sections
 - Single
 - Workshare



Loop Construct

- All programs have loops; For computationally intensive loops; proper for OpenMP parallelisation
- Splits up loop iterations among the threads in a team
- Data parallelism

C/C++:

```
#pragma omp parallel [clauses]
{
    #pragma omp for [clauses] [nowait]
    {
        ...
    }
}
```

Fortran:

```
!$omp parallel [clauses]
    !$omp do [clauses]
        ...
    !$omp end do [nowait]
!$omp end parallel
```

How is OpenMP typically used

- Best sequential program
- Find your most time consuming loops.
- Split them up between threads.

Sequential program

```
void main()
{
    double Res[1000];
    for(int i=0;i<1000;i++)
    {
        do_huge_comp(Res[i]);
    }
}
```

Parallel program

```
void main()
{
    double Res[1000];
    #pragma omp parallel num_threads(4)
    #pragma omp for
    for(int i=0;i<1000;i++)
    {
        do_huge_comp(Res[i]);
    }
}
```

Loop Construct

- Clauses:

`shared (list)`

`private (list)`

`reduction (operator: list)`

`schedule (type [,chunk])`

`ordered`

`firstprivate (list)`

`lastprivate (list)`

`nowait`

Example: (Parallel matrix vector product $y=Ax$)

C - Serial:

```
for(i=0; i<n*n; ++i)
    A[i]=i;

for(i=0; i<n; ++i)
    x[i]=1;

for (i = 0; i < n; ++i)
{
    sum=0;
    for (j=0; j<n; ++j)
        sum+=A[i*n+j]*x[j];
    y[i]=sum;
    printf("y[%d]=%f\n",i, y[i]);
}
```

C - Parallel:

```
#pragma omp parallel
    default(none) ,
    shared(n, A, x, y) ,
    private(i, j, sum, tid)
{
    tid = omp_get_thread_num();
    #pragma omp for
    for (i = 0; i < n; ++i)
    {
        sum=0;
        for (j=0; j<n; ++j)
            sum+=A[i*n+j]*x[j];
        y[i]=sum;
        printf("%d: y[%d]=%f\n", tid,
            i, y[i]);
    }
}
```

Example: (Parallel matrix vector product $y=Ax$)

Fortran - Serial:

```
do j=1,n
  do i=1,n
    A(i,j)=(i-1)*n+j-1
  end do
end do

do i=1,n
  x(i)=1
end do

do i=1,n
  sum=0.0
  do j=1,n
    sum=sum+A(i,j)*x(j)
  end do
  y[i]=sum
  print *, "y[\",i,\""]="\",y[i]
end do
```

Fortran - Parallel:

```
!$omp parallel default(none),
                        shared(n, A, x, y),
                        private(i, j, sum, tid)

  tid=omp_get_thread_num()
  !omp omp do
  do i=1,n
    sum=0.0
    do j=1,n
      sum=sum+A(i,j)*x(j)
    end do
    y[i]=sum
    print *,tid,": y[\",i,\""]="\",y[i]
  end do
  !$omp end do
!$omp end parallel
```




Size n = 16.

Compile: (Intel)

```
$icc hello.c -openmp
```

Execute:

```
$export OMP_NUM_THREADS=4
```

```
$qsub script.pbs
```

Thread 0: y[0]=120.000000

Thread 0: y[1]=376.000000

Thread 0: y[2]=632.000000

Thread 0: y[3]=888.000000

Thread 1: y[4]=1144.000000

Thread 1: y[5]=1400.000000

Thread 1: y[6]=1656.000000

Thread 1: y[7]=1912.000000

Thread 2: y[8]=2168.000000

Thread 2: y[9]=2424.000000

Thread 2: y[10]=2680.000000

Thread 2: y[11]=2936.000000

Thread 3: y[12]=3192.000000

Thread 3: y[13]=3448.000000

Thread 3: y[14]=3704.000000


Thread 3: y[15]=3960.000000

Data dependency


- Two rows may not be updated simultaneously; but columns may be
- For the inner loop: it requires n-1 fork and join
- Invert the loops; better in Fortran

```
for(i=0; i<n; i++)
    for(j=0; j<n; j++)
        a[i][j]=2*a[i-1][j];
```

Loop dependency



```
#pragma omp parallel
for(i=0; i<n; j++)
#pragma omp for
    for(j=0; j<n; j++)
        a[i][j]=2*a[i-1][j];
```



```
#pragma omp parallel
#pragma omp for private(i)
for(j=0; j<n; j++)
    for(i=0; i<n; i++)
        a[i][j]=2*a[i-1][j];
```

Reduction Clause

C - Serial:

```
for(i=0; i<n*n; ++i)
    A[i]=i;

for(i=0; i<n; ++i)
    x[i]=1;

for (i = 0; i < n; ++i)
{
    sum=0;
    for (j=0; j<n; ++j)
        sum+=A[i*n+j]*x[j];
    y[i]=sum;
    printf("y[%d]=%f\n",i, y[i]);
}
```

C - Parallel:

```
#pragma omp parallel
    default(none) ,
    shared(n, A, x, y) ,
    private(i, j, sum, tid)
{
    tid = omp_get_thread_num();
    #pragma omp for
    for (i = 0; i < n; ++i)
    {
        sum=0;
        for (j=0; j<n; ++j)
            sum+=A[i*n+j]*x[j];
        y[i]=sum;
        printf("%d: y[%d]=%f\n", tid,
            i, y[i]);
    }
}
```

Reduction Clause

- Performs a reduction operation on the variables in the list

C/C++: `reduction(operator: list)`

Fortran: `reduction(operator|intrinsic: list)`

- How do we handle this case?

```
sum=0;  
for(i=1; i<n; ++i)  
    sum+=a[i];
```



```
sum=0;  
#pragma omp parallel  
#pragma omp for reduction(+:sum)  
for(i=1; i<n; ++i)  
    sum+=a[i];
```

Reduction Clause

- Performs a reduction operation on the variables in the list

C/C++: `reduction(operator: list)`

Fortran: `reduction(operator|intrinsic: list)`

- How do we handle this case?

```
sum=0
do i=1,n
    sum=sum+a(i)
enddo
```



```
sum=0
!$omp parallel
!$omp do reduction(+:sum)
do i=1,n
    sum=sum+a(i)
enddo
!$omp end do
!$omp end parallel
```

Reduction Clause

- A private copy of each list variable is created for each thread.
- Each thread does the partial reduction.
- At the end, reduction variable is combined with private copies.
- Reduction operations:
 - + , - , * , ... (C/C++/Fortran)
 - && , || , ... (C/C++)
 - .AND. , .OR. , ... (Fortran)
 - max , min , ... (Fortran)

Example: (Parallel matrix vector product $y=Ax$)

C - Serial:

```
for (i=0; i<n*n; ++i)
    A[i]=i;

for (i=0; i<n; ++i)
    x[i]=1;

for (i = 0; i < n; ++i)
{
    sum=0;
    for (j=0; j<n; ++j)
        sum+=A[i*n+j]*x[j];
    y[i]=sum;
    printf("y[%d]=%f\n", i,
y[i]);
}
```

C - Parallel:

```
#pragma omp parallel
    default(none),
    shared(n, A, x, y),
    private(i, j, sum, tid)
{
    tid = omp_get_thread_num();
    #pragma omp for
    for (i = 0; i < n; ++i)
    {
        sum=0;
        for (j=0; j<n; ++j)
            sum+=A[i*n+j]*x[j];
        y[i]=sum;
        printf("%d: y[%d]=%f\n", tid, i,
y[i]);
    }
}
```

Example: (Parallel matrix vector product $y=Ax$)

C - Serial:

```
for (i=0; i<n*n; ++i)
    A[i]=i;

for (i=0; i<n; ++i)
    x[i]=1;

for (i = 0; i < n; ++i)
{
    sum=0;
    for (j=0; j<n; ++j)
        sum+=A[i*n+j]*x[j];
    y[i]=sum;
    printf("y[%d]=%f\n",i, y[i]);
}
```

C:

```
#pragma omp parallel default(shared)
{
for (i = 0; i < n; ++i)
{
    sum=0;
    #pragma omp for private(j, tid)
        reduction(+:sum)
    for (j=0; j<n; ++j)
    {
        sum=sum+(A[i*n+j]*x[j]);
        if(i==0)
        {
            tid = omp_get_thread_num();
            printf("Tid=%d, A[%d]=%f,
sum=%f\n", tid, i*n+j, A[i*n+j], sum);
        }
    }
    y[i]=sum;
    printf("y[%d]=%f\n", i, y[i]);
}
}
```


Example: (Calculation of $y[0]$, $i=0$)

Tid=0, $A[0,0]=0.000000$, sum=0.000000

Tid=0, $A[0,1]=1.000000$, sum=1.000000

Tid=0, $A[0,2]=2.000000$, sum=3.000000

Tid=0, $A[0,3]=3.000000$, sum=6.000000

Tid=1, $A[0,4]=4.000000$, sum=4.000000

Tid=1, $A[0,5]=5.000000$, sum=9.000000

Tid=1, $A[0,6]=6.000000$, sum=15.000000

Tid=1, $A[0,7]=7.000000$, sum=22.000000

Tid=2, $A[0,8]=8.000000$, sum=8.000000

Tid=2, $A[0,9]=9.000000$, sum=17.000000

Tid=2, $A[0,10]=10.000000$, sum=27.000000

Tid=2, $A[0,11]=11.000000$, sum=38.000000

Tid=3, $A[0,12]=12.000000$, sum=12.000000

Tid=3, $A[0,13]=13.000000$, sum=25.000000

Tid=3, $A[0,14]=14.000000$, sum=39.000000

Tid=3, $A[0,15]=15.000000$, sum=54.000000

$y[0]=120.$



Example: (Parallel matrix vector product $y=Ax$)

Fortran:

```

!$omp parallel default(shared)
do i=1,n
  sum=0;
  !$omp do private(j, tid) reduction(+:sum)
  do j=1,n
    sum=sum+A(i,j)*x[j];
    if(i==0) then
      tid=omp_get_thread_num();
      print *, "Tid=", tid, "A[", i, ",", j, "]= ", a(i,j), "sum=", sum
    end if
  end do
  y[i]=sum
  print *, "y[", i, "]= ", y[i]
enddo
!$omp end parallel

```

Loop Construct

- Clauses:

`shared (list)`

`private (list)`

`reduction (operator: list)`

`schedule (type [,chunk])`

`ordered`

`firstprivate (list)`

`lastprivate (list)`

`nowait`

Schedule Clause

- Describes how iterations of the loop are divided among the team threads

C/C++: `schedule (type [,chunk])`

Fortran: `schedule (type [,chunk])`


- Types:

- `schedule (static [,chunk])` : divided into pieces of size chunk, and statically assigned to threads.
- `schedule (dynamic [,chunk])` : divided into pieces of size chunk, and dynamically scheduled as requested.
- `schedule (guided [,chunk])` : size of chunk decreases over time.



Schedule (Static)

- Iterations are divided evenly among threads
- Divides the work into chunk sized parcels
- If there are n threads, each does every n^{th} chunk of work




C/C++:

```
#pragma omp for schedule(static,3)
for (i=0; i<36; ++i)
    Work(i);
```

Schedule (Static)

- Iterations are divided evenly among threads
- Divides the work into chunk sized parcels
- If there are n threads, each does every n^{th} chunk of work



Fortran:


```
!$omp do schedule(static,3)
```

```
do i = 1, 36
  Work (i)
end do
```

```
!$omp end do
```

Schedule(dynamic)

- Threads grab chunks of iterations dynamically
- As a thread finishes one chunk, it grabs the next available chunk
- More overhead than static, but better load balancing



C/C++:

```
#pragma omp for schedule(dynamic,1)
for(i=0;i<36;++i)
    Work(i);
```

Schedule(dynamic)


- Threads grab chunks of iterations dynamically
- As a thread finishes one chunk, it grabs the next available chunk
- More overhead than static, but better load balancing

Fortran:

```
!$omp do schedule(dynamic,1)
```


```
do i = 1, 36
  Work (i)
end do
```

```
!$omp end do
```



Schedule(guided)

- Iterations are divided into chunks such that the size of each successive chunk decreases.
- chunk: the size of the smallest chunk size
- Less overhead than dynm., good l. b.




C/C++:

```
#pragma omp for schedule(guided,1)
for(i=0;i<36;++i)
    Work(i);
```

Schedule(guided)

- Iterations are divided into chunks such that the size of each successive chunk decreases.
- chunk: the size of the smallest chunk size
- Less overhead than dynm., good l. b.



Fortran:

```
!$omp do schedule(guided,1)
```

```
do i = 1, 36
  Work (i)
end do
```

```
!$omp end do
```

Work-Sharing Constructs

- Work-sharing construct divides work among the thread team; what kind of work to be parallelised.
- Inside the parallel region.
- No new threads created. Construct must be encountered by all threads in the team.
- No implied barrier on entry to a work-sharing construct; Yes at end of construct.
- Work-sharing Constructs:
 - Loop
 - Sections
 - Single
 - Workshare

Sections Construct

- Specifies enclosed sections that are distributed over the threads in the team; implicit barrier at the end
- Each section is executed by one thread.
- Functional parallelism

C/C++:

```
#pragma omp parallel [clauses]
{
  #pragma omp sections [clauses] nowait
  {
    #pragma omp section
    ...
    #pragma omp section
    ...
  }
}
```


Fortran:

```
!$omp parallel [clauses]
!$omp sections [clauses]
  !$omp section
  ...
  !$omp section
  ...
!$omp end sections
[nowait]
!$omp end parallel
```

```
$export OMP_NUM_THREADS=4
```

```
Hello from thread 0
Hello from thread 2
Hello from thread 3
```

```
#include <stdio.h>
#include <omp.h>
int main()
{
    int tid;
    #pragma omp parallel private(tid)
    {
        tid = omp_get_thread_num();
        #pragma omp sections
        {
            #pragma omp section
            printf("Hello from thread %d \n", tid);
            #pragma omp section
            printf("Hello from thread %d \n", tid);
            #pragma omp section
            printf("Hello from thread %d \n", tid);
        }
    }
}
```



Sections Construct

- Clauses:

`private (list)`

`firstprivate (list)`

`lastprivate (list)`

`reduction (operator: list)`

`nowait`

Work-Sharing Constructs

- Work-sharing construct divides work among the thread team; what kind of work to be parallelised.
- Inside the parallel region.
- No new threads created. Construct must be encountered by all threads in the team.
- No implied barrier on entry to a work-sharing construct; Yes at end of construct.
- Work-sharing Constructs:
 - Loop
 - Sections
 - Single
 - Workshare

Single Construct

- Specifies a block of code that is executed by only one of the threads in the team.
- Rest of threads wait at the end of enclosed code block
- May be useful when dealing with sections of code that are not thread-safe.

C/C++:

```
#pragma omp parallel [clauses]
{
    #pragma omp single [clauses]
    ...
}
```

Fortran:

```
!$omp parallel [clauses]
    !$omp single [clauses]
    ...
    !$omp end single
!$omp end parallel
```


Single Construct

- **Clauses:**

```
private (list)
firstprivate (list)
copyprivate (list)
nowait
```

```
#pragma omp parallel num_threads(4)
{
    #pragma omp single copyprivate(a)
    read_from_file(a);

    compute(a);

    #pragma omp single
    write_to_file(result);
}
```

- **Copyprivate(*list*):** used to broadcast private variable values from a single thread to all instances of the private variables in the other threads.

Work-Sharing Constructs

- Work-sharing construct divides work among the thread team; what kind of work to be parallelised.
- Inside the parallel region.
- No new threads created. Construct must be encountered by all threads in the team.
- No implied barrier on entry to a work-sharing construct; Yes at end of construct.
- Work-sharing Constructs:
 - Loop
 - Sections
 - Single
 - **Workshare**

Workshare construct

- Fortran only
- Divides the execution of the enclosed structured block into separate units of work
- Threads of the team share the work
- Each unit is executed only once by one thread
- Allows parallelisation of
 - array and scalar assignments
 - WHERE statements and constructs
 - FORALL statements and constructs
 - parallel, atomic, critical constructs

```
!$omp workshare  
...  
!$omp end workshare  
[nowait]
```

```

program wshare
use omp_lib
implicit none

integer :: i
real :: a(10), b(10), c(10)
do i=1,10
    a(i)=i
    b(i)=i+1
enddo

!$omp parallel shared(a, b, c)
!$omp workshare
    c=a+b
!$omp end workshare nowait
!$omp end parallel

end program wshare
  
```

Combined Constructs

- The following shortcuts are supported:
 - parallel for / parallel do
 - parallel sections
 - parallel workshare
- Equivalent to a parallel construct followed by a work-sharing construct except nowait.

```
#pragma omp parallel for
```

Same as

```
#pragma omp parallel
#pragma omp for
```

```
!$omp parallel do
```

Same as

```
!$omp parallel
!$omp do
```

Parallel Loop Directives

C/C++:

```
#pragma omp parallel for [clause[clause]...]
for ( index = first; index <= last ; index++ )
{
    body of the loop
}
```

Fortran:

```
!$omp parallel do [clause[clause]...]
do index = first, last [, stride]
    body of the loop
enddo
!$omp end parallel do
```

- Readability
- Performance advantage

Outline

- Introduction to OpenMP
- OpenMP Directives
 - Directive Format
 - Parallel Construct
 - Work-Sharing Constructs
 - Synchronization Constructs
- Performance Considerations
- OpenMP Versions and Features

Synchronization Constructs

- Threads communicate through shared variables.
- Uncoordinated access of these variables can lead to undesired effects.
 1. Two threads update (write) a shared variable in the same step of execution, the result is dependent on the way this variable is accessed. This is called a race condition.
 2. Suppose that one processor has an updated result in private cache. Second processor wants to access that memory location - but a read from memory will get the old value since original data not yet written back.



Synchronization Constructs

- Synchronization imposes order constraints.
- Used to protect access to shared data.
- High-Level Synchronization Constructs:
 - master
 - critical
 - barrier
 - atomic
 - ordered
- Low-Level Synchronization Constructs:
 - flush
 - locks

Master directive:

- Only master thread can enter the structured block

C/C++:

```
#pragma omp master
{
    ...
}
```

Fortran:

```
!$omp master
...
!$omp end master
```

Critical directive:

- Only one thread at a time can enter a critical section

```
#pragma omp critical [name]
{
    ...
}
```

```
!$omp critical [name]
...
!$omp end critical
```

```
#pragma omp parallel
```

```
.....
```

```
#pragma omp critical ←
    result=result+temp;
```

Threads wait here: only one thread at a time executes the statement. So this is a piece of sequential code.

Barrier directive:

- Threads wait each other until all have reached that barrier

C/C++:

```
#pragma omp barrier
```

```
for(i= 0; i<N; i++)
    a[i] = b[i] + c[i];
#pragma omp barrier ←
for(i= 0; i<N; i++)
    d[i] = a[i] + b[i];
```

Fortran:

```
!$omp barrier
```

Threads wait here and only continue when all threads have reached the barrier point.

Atomic directive:

- Single thread can access a memory location at a time.
- $x \text{ binop} = \text{expr}$, $x++$, $--x$, ... (C/C++)
- $x = x \text{ op } \text{expr}$, $x = \text{expr op } x$, ... (Fortran)

```
#pragma omp parallel
```

```
.....
```

```
#pragma omp atomic
```

```
    x+=temp;
```

C/C++:

```
#pragma omp atomic  
    statement
```

Fortran:

```
!$omp atomic  
    statement
```

Ordered directive:

- Iterations of the enclosed loop will be executed in the same order as if they were executed sequentially

C/C++:

```
#pragma omp for ordered
...
#pragma ordered ...
```

```
#pragma omp parallel for ordered
for(i=0;i<N;i++)
{
    #pragma ordered ←
    printf(" %d\n", i);
}
```

Fortran:

```
!$omp do ordered
...
!$omp ordered
...
!$omp end ordered
!$omp end do
```

The output always will be
0, 1, 2, ..., N

Synchronization Constructs

- Synchronization imposes order constraints and is used to protect access to shared data.
- High-Level Synchronization Constructs:
 - master
 - critical
 - barrier
 - atomic
 - ordered
- Low-Level Synchronization Constructs:
 - flush
 - locks

Flush directive:

- A synchronization point at which thread visible variables are written back to the memory

C/C++:

```
#pragma omp flush (list)
```

Fortran:

```
!$omp flush (list)
```

```
#pragma omp section
```

```
a=compute () ;
```

```
flush(a) ;
```

Locks

- Occasionally we may require more flexibility than is provided by critical and atomic directives.
- A **lock** is a special variable that may be **set** by a thread. No other thread may **set** the lock until the thread which set the lock has **unset** it.
- A lock must be initialised before it is used, and may be destroyed when it is no longer required.
- Lock variables should not be used for any other purpose.

C/C++:

```
omp_lock_t lock;
```

```
void omp_init_lock(omp_lock_t *lock)
void omp_destroy_lock(omp_lock_t *lock)
void omp_set_lock(omp_lock_t *lock)
void omp_unset_lock(omp_lock_t *lock)
int omp_test_lock(omp_lock_t *lock)
```

Fortran:

```
integer(omp_lock_kind) :: lock;
```

```
subroutine omp_init_lock(lock);
subroutine omp_destroy_lock(lock);
subroutine omp_set_lock(lock);
subroutine omp_unset_lock(lock);
subroutine omp_test_lock(lock);
```

- There are also nestable lock routines which allow the same thread to set a lock multiple times before unsetting it the same number of times.

Lock Example

```
omp_lock_t writelock;  
omp_init_lock(&writelock);  
#pragma omp parallel for  
for ( i = 0; i < x; i++ )  
{  
    // some stuff  
    omp_set_lock(&writelock);  
    // one thread at a time stuff  
    omp_unset_lock(&writelock);  
    // some stuff  
}  
  
omp_destroy_lock(&writelock);
```

Timing with OpenMP

- **omp_get_wtime**

C/C++: `double omp_get_wtime(void)`

Fortran: `double precision function omp_get_wtime()`

- Returns a double-precision floating point value equal to the number of elapsed seconds since some point in the past.

```
double t;
t = omp_get_wtime();
... work to be timed ...
t = omp_get_wtime() - t;
```

```
real(8) :: t
t= omp_get_wtime()
... work to be timed ...
t= omp_get_wtime()-t
```

- **omp_get_wtick:** Returns the number of seconds between processor clock ticks.

Outline

- Introduction to OpenMP
- OpenMP Directives
 - Directive format
 - Parallel Construct
 - Work-Sharing Constructs
 - Synchronization Constructs
- Performance Considerations
- OpenMP versions and features

Correctness vs Performance

- It may be easy to write a correctly functioning OpenMP program.
- But, it is not so easy to create a program that provides the desired level of performance!

Basic Strategies

- If a parallelized loop does not perform well, consider
 - Thread start up costs, The amount of time spent handling OpenMP constructs, Thread termination time
 - Avoid parallel overhead at low iteration counts

```
#pragma omp parallel for if(M > 800)
for(j=0; j< M; j++)
    aa[j] = alpha*bb[j] + cc[j];
```

- Load imbalances: Unequal work loads lead to idle threads and wasted time. Use proper scheduling type.
- Unnecessary synchronization; Large Critical Regions; prefer atomic update if possible


Basic Strategies

- Many references to shared variables
 - Use private data - allocated on stack
- Low cache reuse
 - If each thread accesses a distinct portion of data consistently through the program, then they will load this data to their local cache and they don't need to go to memory for each reference.
 - In C, a 2-D array is stored in rows (and in columns in Fortran). Organize data accesses so that values are used as often as possible while they are still in the cache.
- False sharing
 - when two threads update different data elements in the same cache line.



This invalidates the cache line and forces a memory update to maintain cache coherency.

It can be reduced by making use of private data as much as possible



Debugging OpenMP Code

- Shared memory parallel programming opens up a range of new programming errors arising from unanticipated conflicts between shared resources.
- Race conditions:
 - Multiple threads are updating the same shared variable simultaneously.
 - Hard to find, not reproducible, answer varies with number of threads.
- Deadlock:
 - When threads hang while waiting on a locked resource that will never become available.
 - A simple approach is to put print statement in front of all lock calls.

Deadlock

```

omp_lock_t  lock;
omp_init_lock(&lock);
#pragma omp parallel sections
{
#pragma omp section
{
    omp_set_lock(&lock)
    ierr=work1();
    if (ierr==0) {
        omp_unset_lock(&lock);
    }else{
        printf("Error\n");
    }
}
#pragma omp section
{
    omp_set_lock(&lock);
    ierr=work2();
    omp_unset_lock(&lock);
}
}
omp_destroy_lock(&lock);
  
```

Deadlock

```
integer(omp_kind_lock) :: ilock
call omp_init_lock(ilock)
!$omp parallel sections
!$omp section
    call omp_set_lock(ilock)
    call work1(ierr)
    if (ierr.eq.0) then
        call omp_unset_lock(ilock)
    else
        print *, "Error"
    endif
!$omp section
    call omp_set_lock(ilock)
    call work2(ierr)
    call omp_unset_lock(ilock)
!$omp end parallel sections

call omp_destroy_lock(ilock)
```

Using Print Statements

- Advantages:

- simple
- useful for deterministic bugs
- monitoring the iterations on threads

- Disadvantages:

- slow
- specification of what to display
- human intensive bug hunting

Debugging and Profiling Tools

- Debugging Tools:
 - Intel debuggers: gdb-ia
 - Intel Inspector
 - DDT
 - Totalview
- Profiling Tools:
 - Valgrind
 - GNU profiler: gprof
 - Intel VTune
 - ompP
 - Tau
 - VampirTrace




General Performance Considerations

- Be aware of directives cost
 - Parallelize outer loops
 - Minimize the number of directives
 - Minimize synchronization - minimize the use of barrier, critical, ordered
 - Consider using nowait clause of omp for/do when enclosing several loops inside one parallel region.
 - Merge loops to reduce synchronization cost
 - Maximize parallel regions

- Be aware of the Amdahl's law
 - Minimize serial code
 - Remove dependencies among iterations
- Balance the load
 - Experiment with using schedule clause
- Reduce false sharing
 - Use private variables
- Try task level parallelism



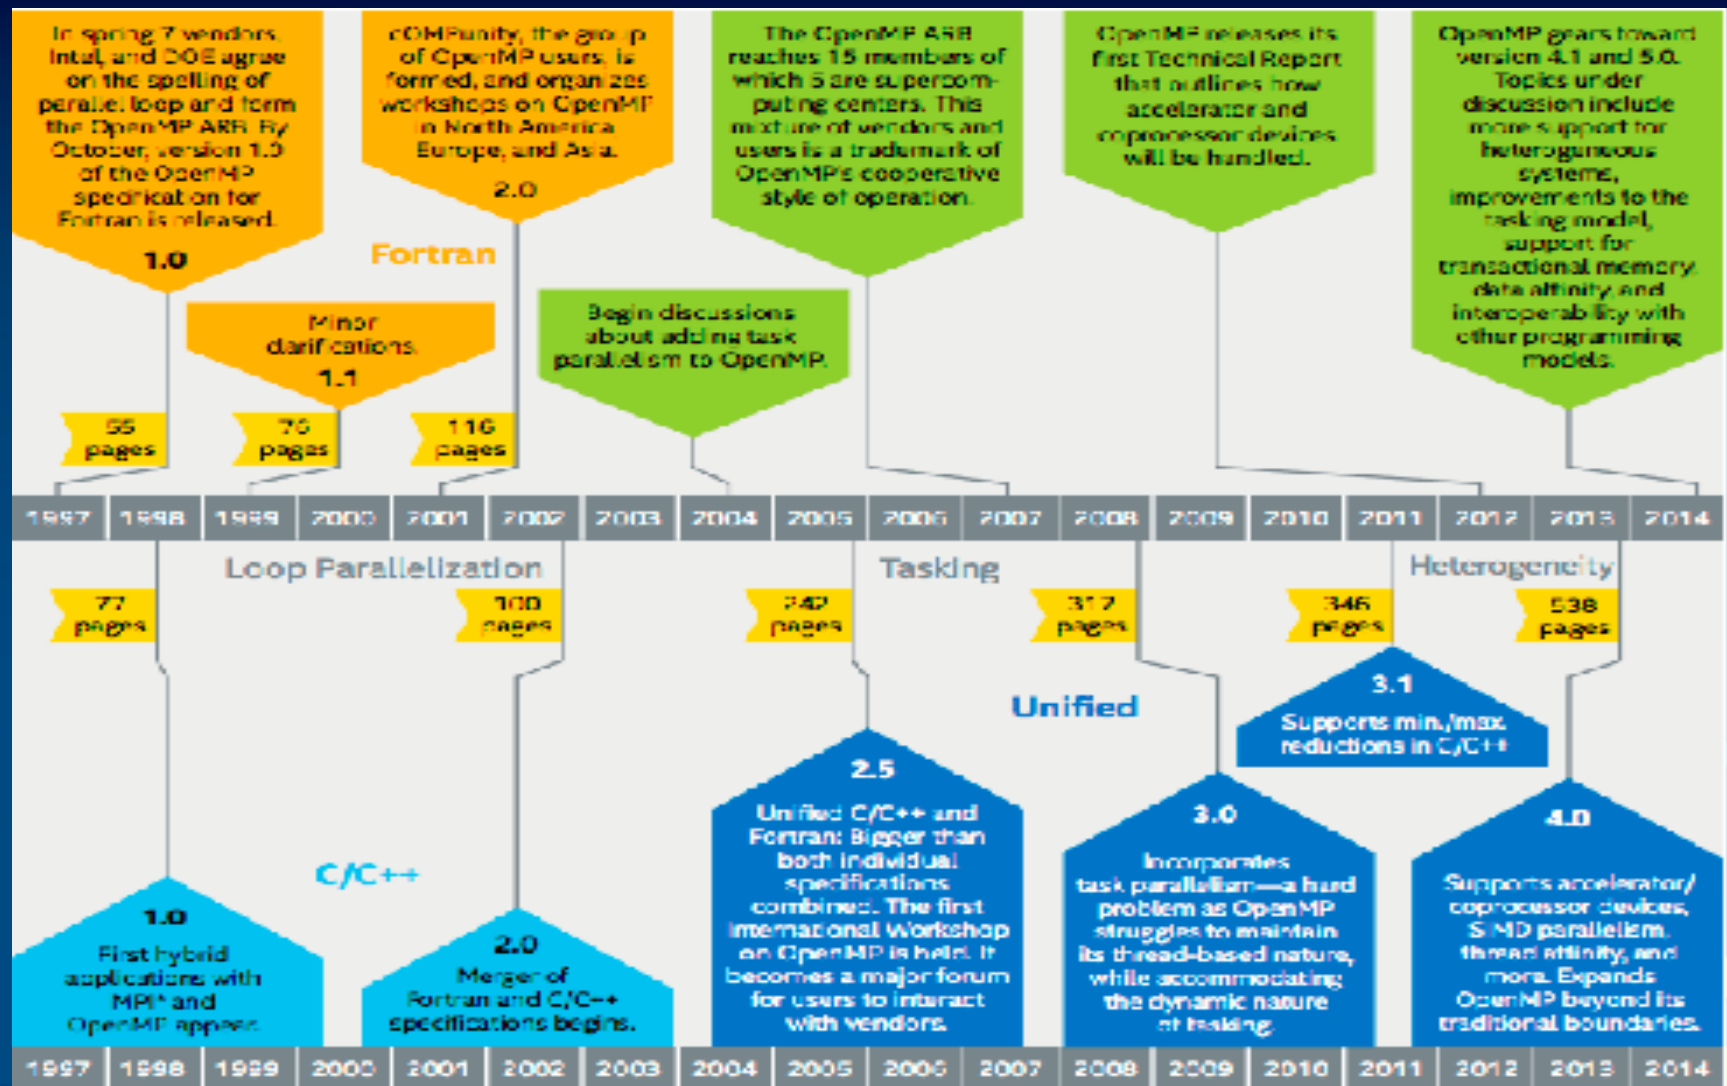
OpenMP Directives Overhead



Outline

- Introduction to OpenMP
- OpenMP Directives
 - Directive format
 - Parallel Construct
 - Work-Sharing Constructs
 - Synchronization Constructs
- Performance Considerations
- OpenMP versions and features

OpenMP Timeline



OpenMP 3.0

Version 3.0 released in May 2008

- New task level parallelism
- Improvements to loop and nested parallelism
- Additional Clauses, runtime functions and environment variables

Version 3.0 Features

- Adding tasking is the biggest addition for 3.0
- OpenMP has always had tasks, never called them explicitly that.
- For parallelizing irregular problems, unbounded loops, recursive algorithms, multi-block grids and many others

C/C++:

```
#pragma omp task [clauses]
{
    ...
}
```

Fortran:

```
!$omp task [clauses]
...
!$omp end task
```

- Clauses: if, untied, shared, private, firstprivate, default
 - When if is false, executed immediately by the encountering thread.
 - Tied by default: The same thread from beginning to end will execute the code

- untied: to specify that different threads execute different parts of the code
- Tasks can create descendants to form a task tree. They may run in parallel to the parent or suspend and run synchronously.
- Task Synchronisation:
 - taskwait: Waits for all child tasks to be completed.
 - C/C++:** `#pragma omp taskwait`
 - Fortran:** `!$omp taskwait`
 - Can be controlled by `omp barrier` / `implicit barrier`

```
#pragma omp parallel num_threads(n)
{
    #pragma omp task
    function_A();
    #pragma omp barrier
    #pragma omp single
    {
        #pragma omp task
        function_B();
    }
}
```

Ex: (Fibonacci Sequence)

$$\text{fib}(0)=0$$

$$\text{fib}(1)=1$$

$$\text{fib}(n)=\text{fib}(n-1)+\text{fib}(n-2), n>1$$

0, 1, 1, 2, 3, 5, 8, 13, 21, 34, ...



```
int main(){
    int n=50;

    omp_set_num_threads(4);

    #pragma omp parallel shared(n)
    {
        #pragma omp single
        printf("fib(%d)=%d\n", n, fib(n));
    }
}
```

```
int fib(int n)
{
    int i, j;
    if(n<2) return n;
    else{
        #pragma omp task shared(i) firstprivate(n)
        i=fib(n-1);
        #pragma omp task shared(j) firstprivate(n)
        j=fib(n-2);
        #pragma omp taskwait
        return i+j;
    }
}
```

- Better support for nested parallelism

- controls the maximum number of nested active parallel regions

`omp_set_max_active_levels()`, `OMP_MAX_ACTIVE_LEVELS`

- sets the number of OpenMP threads to use for the whole OpenMP program

`omp_get_thread_limit()`, `OMP_THREAD_LIMIT`

- Improvements to loop parallelism

- `schedule(runtime)`: to set iterations at runtime through the environment variable `OMP_SCHEDULE`
- `schedule(auto)`: The compiler or runtime system decides what is best to use depending on the implementation

collapse(n) clause: parallelize perfectly nested multi-dimensional loops; takes a constant positive integer as a parameter, which determines how many loops are collapsed. Compiler then forms a single loop and parallelizes it.

```
#pragma omp parallel for collapse(2)
for(i=0;i<n;i++)
    for(j=0;j<m;j++)
        ...
```

OMP_STACK_SIZE size [B|K|M|G]:
control of thread's stack size

OpenMP 3.1

Version 3.1 released in July 2011

- Additional Clauses
- Improvements to task parallelism
- Initial support for thread binding



Version 3.1 Features

- Extensions to OpenMP tasking model
 - Taskyield construct: define task-switching points; suspend the current task and execute another
- C/C++: `#pragma omp taskyield`
- Fortran: `!$omp taskyield`
- **mergeable** clauses: a task can have the same data region as the generating task region; avoid potentially expensive initialization of the task environment.
 - **final** clause: a task that makes all its child tasks to be executed sequentially in the same region; a task may not be scheduled for deferred execution, but instead is immediately executed.

Version 3.1 Features

- `omp_in_final()`: returns true if the calling task is final.
- Initial support for thread binding: Control whether OpenMP threads are allowed to move between processors or not
 - `OMP_PROC_BIND` to TRUE for not moving threads
- Control nested thread team sizes: export `OMP_NUM_THREADS=n1, n2, n3`
- For C/C++, the reduction clause now accepts min and max functions.
- The atomic construct now accepts the clauses: read, write, update and capture

OpenMP 4.0

OpenMP 4.0 released in July 2013

- SIMD directives
- Extended support for thread affinity
- New user-defined reductions
- Error Handling
- Accelerators support
- Support for Fortran 2003

Version 4.0 Features

- Simd Construct: Transform the loop into a simd loop
- multiple iterations of the loop can be executed concurrently
- Loop to be executed using SIMD lanes:

C/C++:

```
#pragma omp simd [clauses]  
  for – loops
```

Fortran:

```
!$omp simd [clauses]  
  do – loops  
!$omp end simd
```

- Function that can be called from a SIMD loop.

```
#pragma omp declare simd [clauses]  
  function
```

```
!$omp declare simd [clauses]  
  function  
!$omp end simd
```

Version 4.0 Features

Clauses for simd: safelen, linear, aligned, private, lastprivate, reduction, collapse

- safelen (length): limits the number of iterations in a SIMD chunk
- linear (list): declares a number of list items to be private to a SIMD lane
- aligned (list): declares a number of items to be aligned to some number of bytes

Clauses for declare simd: simdlen, linear, aligned, uniform, reduction, inbranch, notinbranch

- simdlen (length): number of concurrent arguments for the function
- uniform (argument): arguments to have an invariant value for all function invocations
- inbranch/notinbranch: the function will always/never be called from inside an if statement of a SIMD loop



- New Reduction Clause: use your own reduction operation on your own type

C/C++:

```
#pragma omp declare reduction(reduction-identifier : typename-list : combiner) [initializer-clause]
```

Fortran:

```
!$omp declare reduction(reduction-identifier : type-list : combiner) [initializer-clause]
```

- Reduction-identifier: gives a name to the operator
- Typename-list: A list of types to which it applies
- Combiner: specifies how to combine values
- Initializer-clause: specifies how to initialise the private elements of each thread


```
#pragma omp declare reduction (merge : std::vector<int> :  
omp_out.insert(omp_out.end(), omp_in.begin(), omp_in.end()))
```

```
void schedule (std::vector<int> &v, std::vector<int> &filtered) {  
#pragma omp parallel for reduction (merge : filtered)  
    for (std::vector<int>::iterator it = v.begin(); it < v.end(); it++)  
        if ( filter(*it)) filtered.push_back(*it);  
}
```

- omp_out refers to private copy that holds combined value
- omp_in refers to the other private copy

- Affinity Support: better locality, less false sharing, more memory bandwidth.
 - New Clause to the parallel construct: specifies the places to use for the threads in the team
`proc_bind(master|close|spread)`
master: threads to the same place as the master thread
close: threads close to the place of the master thread
spread: spread threads across the machine
 - `OMP_PROC_BIND` can now specify master, close, spread
 - New RTLs: `omp_proc_bind_t omp_get_proc_bind(void)`
 - `OMP_PLACES`: bind the OpenMP threads to the places in the list in terms of threads, core, sockets

- Error Model:

- improve the stability of OpenMP applications against system-level, runtime-level, and user-defined errors.
- based on conditional cancellation and user-defined cancellation points
- Cancel construct: cancellation of all tasks in the same construct

C/C++:

```
#pragma omp cancel [clauses]
```

Fortran:

```
!$omp cancel [clauses]
```

Clauses: parallel, sections, for/do, taskgroup, if

taskgroup: marks a region such that all tasks started in it belong to a group

- **OMP_CANCELLATION** set to TRUE to enable

- Accelerator model:

- Enables the usage of accelerators and coprocessors to offload computation
- target construct: marks a region to execute on device
 - Creates device tasks that are executed by device-only threads.
- A device has a data environment. Variables are copied in and out of a data environment or

C/C++:

```
#pragma omp target [clauses]
```

```
...
```


Fortran:

```
!$omp target [clauses]
```

```
...
```

```
!$omp end target
```

- Clauses: if, device(expression), map ([map-type:] list)



- `to`: existing host variables copied to a corresponding variable in the target before
- `from`: target variables copied back to a corresponding variable in the host after
- `tofrom`: Both from and to
- `alloc`: Neither from nor to, but ensure the variable exists on the target but no relation to host variable.

OpenMP 5.0 Plans

- Support for memory affinity
- Refinements to accelerator support
- Additional task/thread synchronization mechanisms
- Completing extension of OpenMP to Fortran 2003

Conclusions

- OpenMP is a parallel programming model for SMP machines.
- All threads have access to a shared main memory; each thread can have local memory.
- The parallelism has to be decided explicitly by the programmer: Data parallelism, Task parallelism
- To control the parallelization, thread exclusion and synchronization constructs can be used.



- OpenMP is successful in small-to-medium SMP systems.
- Multiple cores/CPU's dominate the future computer architectures; OpenMP would be the major parallel programming language in these architectures.
- Simple: everybody can learn it in 2 weeks.
- Not so simple: Don't stop learning! Keep learning it for better performance.



References

- <http://openmp.org>
- <https://computing.llnl.gov/tutorials/openMP>
- <http://openmp.org/wp/openmp-specifications>
- Michael J. Quinn, Parallel Programming in C with MPI and OpenMP, Mc Graw Hill, 2003.
- Barbara Chapman, Gabriele Jost and Ruud Van Der Pas, Using OpenMP: Portable Shared Memory Parallel Programming, Volume 10, MIT Press, 2008.